

Name: Mark Abrenio  
Project: Udacity Project 1

For this project, we had to solve the Banana Navigation environment. My algorithm used to solve this challenge was very simple. It is just a standard DQN that is used with a very low GAMMA / Discount rate. In fact, my discount rate was zero!

```
BUFFER_SIZE = int(1e5) # replay buffer size
BATCH_SIZE = 128      # minibatch size
#GAMMA = 0.99          # discount factor
GAMMA = 0.0
#TAU = 1e-3            # for soft update of target parameters
TAU = 1
#LR = 5e-4             # learning rate
LR = 0.0005
UPDATE_EVERY = 4       # how often to update the network
UPDATE_TARGET_EVERY = 10000
|
device = torch.device("cuda:0" if torch.cuda.is_available() else "cpu")
```

I started out trying to use a higher GAMMA value but that didn't work too well because the agent kept taking non optimal paths and it would run out of time. I also used a very small model for this DQN with a small amount of parameters.

```

import torch
import torch.nn as nn
import torch.nn.functional as F

class QNetwork(nn.Module):
    """Actor (Policy) Model."""

    def __init__(self, state_size, action_size, seed, fc1_unit=32, fc2_unit=32):
        """Initialize parameters and build model.

        Params
        =====
            state_size (int): Dimension of each state
            action_size (int): Dimension of each action
            seed (int): Random seed
        """
        super(QNetwork, self).__init__()
        self.seed = torch.manual_seed(seed)
        "*** YOUR CODE HERE ***"
        self.fc1 = nn.Linear(state_size, fc1_unit)
        self.fc2 = nn.Linear(fc1_unit, fc2_unit)
        self.fc3 = nn.Linear(fc2_unit, action_size)

    def forward(self, state):
        """Build a network that maps state -> action values."""
        x = F.relu(self.fc1(state))
        x = F.relu(self.fc2(x))
        return self.fc3(x)

```

As illustrated its only got 3 linear layers. And the fc1 and fc2 layers only output 32 units each. There is really nothing very interesting about the model. I originally tried deeper models and they worked but it takes longer to train. And since the state space isn't really that big it just didn't seem necessary to have all of those parameters! Then I had to set the learning agent to only update the target network every 10000 steps. If the number of steps before updates was small, the agent never really got that good.

```

LR = 0.0005
UPDATE_EVERY = 4          # how often to update the network
UPDATE_TARGET_EVERY = 10000

device = torch.device("cuda:0" if torch.cuda.is_available() else "cpu")

class Agent():
    """Interacts with and learns from the environment."""

    def __init__(self, state_size, action_size, seed):
        """Initialize an Agent object.

        Params
        =====
            state_size (int): dimension of each state
            action_size (int): dimension of each action
            seed (int): random seed
        """
        self.state_size = state_size
        self.action_size = action_size
        self.seed = random.seed(seed)
        self.criterion = torch.nn.MSELoss()

        # Q-Network
        self.qnetwork_local = QNetwork(state_size, action_size, seed).to(device)
        self.qnetwork_target = QNetwork(state_size, action_size, seed).to(device)
        self.optimizer = optim.Adam(self.qnetwork_local.parameters(), lr=LR)

        # Replay memory
        self.memory = ReplayBuffer(action_size, BUFFER_SIZE, BATCH_SIZE, seed)
        # Initialize time step (for updating every UPDATE_EVERY steps)
        self.t_step = 10
        self.l_step = 1

    def STEP() SIMPLY RECORDS AND LEARNS

```

```

def learn(self, experiences, gamma): ###learns and updates
    """Update value parameters using given batch of experience tuples.

    Params
    =====
    experiences (Tuple[torch.Variable]): tuple of (s, a, r, s', done) tuples
    gamma (float): discount factor
    """

    states, actions, rewards, next_states, dones = experiences

    ## TODO: compute and minimize the loss
    *** YOUR CODE HERE ***
    self.qnetwork_local.train()
    self.qnetwork_target.eval()
    predicted_targets = self.qnetwork_local(states).gather(1, actions)
    with torch.no_grad():
        #target_net_pred1 = self.qnetwork_target(next_states)
        target_net_pred = self.qnetwork_target(next_states).max(1)[0].unsqueeze(1) ## **TARGET NETWORK TO GET `MAXQ`


    labels = rewards + (gamma * target_net_pred * (1 - done)s))

    loss = self.criterion(predicted_targets, labels).to(device)
    self.optimizer.zero_grad()
    loss.backward()
    self.optimizer.step()

    # ----- update target network ----- #
    # **by default updates target network every time we learn. but lets
    # **try to only update target every so often
    self.l_step = (self.l_step + 1) % UPDATE_TARGET_EVERY
    #print("learning")
    if self.l_step == 0:
        #print("updating target model")
        self.soft_update(self.qnetwork_local, self.qnetwork_target, TAU)

```

Note: It took the model a couple hundred games to learn the environment. But the code got interrupted before I could let it all train. Therefore in my submission the game only runs approximately 200 times. The first 100 steps it ran with a 90 percent epsilon rate so it could just find some more details about the environment. Then the last 100 step it was playing with epsilon of zero. Below is the chart of the training. As illustrated the first half doesn't go to well because epsilon is 90 percent and the model is still learning some stuff. But the last 100 games go much better because epsilon is turned to zero percent.



Also please note, the final model is saved as the checkpoint file  
`MODEL_CHECKPOINT.2004788.model.pt`. This file can be found no the GitHub repo.

```
In [3]: print("Saving model")
print("Graph of Scores:")
my_agent.save_model() ##The required checkpoint model
from matplotlib import pyplot as plt
plt.plot(all_scores) ##We'll do a little score plot
plt.show()
# When finished, you can close the environment.


# In[ ]:
```

```
env2.close()
```

```
Saving model
```

```
Graph of Scores:
```

```
Model Saved: MODEL_CHECKPOINT.2004788.model.pt
```



The final 100 scores are displayed below. These

scores are reported on the Jupiter notebook as well.

GAME NUMBER:	101
Score:	9.0
GAME NUMBER:	102
Score:	21.0
GAME NUMBER:	103
Score:	16.0
GAME NUMBER:	104
Score:	15.0
GAME NUMBER:	105
Score:	11.0
GAME NUMBER:	106
Score:	24.0
GAME NUMBER:	107
Score:	15.0
GAME NUMBER:	108
Score:	17.0
GAME NUMBER:	109
Score:	4.0
GAME NUMBER:	110
Score:	18.0
GAME NUMBER:	111
Score:	14.0
GAME NUMBER:	112
Score:	18.0
GAME NUMBER:	113
Score:	18.0
GAME NUMBER:	114
Score:	11.0
GAME NUMBER:	115
Score:	17.0
GAME NUMBER:	116
Score:	18.0
GAME NUMBER:	117
Score:	14.0
GAME NUMBER:	118
Score:	16.0
GAME NUMBER:	119
Score:	18.0
GAME NUMBER:	120

Score: 13.0  
GAME NUMBER: 121  
Score: 12.0  
GAME NUMBER: 122  
Score: 20.0  
GAME NUMBER: 123  
Score: 20.0  
GAME NUMBER: 124  
Score: 10.0  
GAME NUMBER: 125  
Score: 14.0  
GAME NUMBER: 126  
Score: 15.0  
GAME NUMBER: 127  
Score: 10.0  
GAME NUMBER: 128  
Score: 19.0  
GAME NUMBER: 129  
Score: 13.0  
GAME NUMBER: 130  
Score: 16.0  
GAME NUMBER: 131  
Score: 17.0  
GAME NUMBER: 132  
Score: 19.0  
GAME NUMBER: 133  
Score: 14.0  
GAME NUMBER: 134  
Score: 15.0  
GAME NUMBER: 135  
Score: 19.0  
GAME NUMBER: 136  
Score: 18.0  
GAME NUMBER: 137  
Score: 24.0  
GAME NUMBER: 138  
Score: 17.0  
GAME NUMBER: 139  
Score: 17.0  
GAME NUMBER: 140  
Score: 15.0

GAME NUMBER: 141  
Score: 13.0  
GAME NUMBER: 142  
Score: 13.0  
GAME NUMBER: 143  
Score: 12.0  
GAME NUMBER: 144  
Score: 19.0  
GAME NUMBER: 145  
Score: 19.0  
GAME NUMBER: 146  
Score: 16.0  
GAME NUMBER: 147  
Score: 13.0  
GAME NUMBER: 148  
Score: 15.0  
GAME NUMBER: 149  
Score: 14.0  
GAME NUMBER: 150  
Score: 13.0  
GAME NUMBER: 151  
Score: 16.0  
GAME NUMBER: 152  
Score: 13.0  
GAME NUMBER: 153  
Score: 14.0  
GAME NUMBER: 154  
Score: 18.0  
GAME NUMBER: 155  
Score: 22.0  
GAME NUMBER: 156  
Score: 13.0  
GAME NUMBER: 157  
Score: 14.0  
GAME NUMBER: 158  
Score: 16.0  
GAME NUMBER: 159  
Score: 18.0  
GAME NUMBER: 160  
Score: 21.0  
GAME NUMBER: 161



Score: 18.0  
GAME NUMBER: 162  
Score: 20.0  
GAME NUMBER: 163  
Score: 10.0  
GAME NUMBER: 164  
Score: 17.0  
GAME NUMBER: 165  
Score: 17.0  
GAME NUMBER: 166  
Score: 18.0  
GAME NUMBER: 167  
Score: 12.0  
GAME NUMBER: 168  
Score: 16.0  
GAME NUMBER: 169  
Score: 22.0  
GAME NUMBER: 170  
Score: 23.0  
GAME NUMBER: 171  
Score: 22.0  
GAME NUMBER: 172  
Score: 16.0  
GAME NUMBER: 173  
Score: 11.0  
GAME NUMBER: 174  
Score: 16.0  
GAME NUMBER: 175  
Score: 12.0  
GAME NUMBER: 176  
Score: 15.0  
GAME NUMBER: 177  
Score: 17.0  
GAME NUMBER: 178  
Score: 16.0  
GAME NUMBER: 179  
Score: 19.0  
GAME NUMBER: 180  
Score: 15.0  
GAME NUMBER: 181  
Score: 15.0

GAME NUMBER: 182

Score: 21.0

GAME NUMBER: 183

Score: 13.0

GAME NUMBER: 184

Score: 11.0

GAME NUMBER: 185

Score: 17.0

GAME NUMBER: 186

Score: 16.0

GAME NUMBER: 187

Score: 14.0

GAME NUMBER: 188

Score: 15.0

GAME NUMBER: 189

Score: 18.0

GAME NUMBER: 190

Score: 25.0

GAME NUMBER: 191

Score: 20.0

GAME NUMBER: 192

Score: 10.0

GAME NUMBER: 193

Score: 10.0

GAME NUMBER: 194

Score: 18.0

GAME NUMBER: 195

Score: 5.0

GAME NUMBER: 196

Score: 22.0

GAME NUMBER: 197

Score: 19.0

GAME NUMBER: 198

Score: 16.0

GAME NUMBER: 199

Score: 16.0

GAME NUMBER: 200

WRITING MODEL.

Model Saved: MODEL\_CHECKPOINT.6652469.model.pt

Scores 13 or above: 84

Scores below 13: 16

Average Score over Last 100: 15.94